



Project By : PRASAD JADHAV

```
In [1]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

from imblearn.over_sampling import SMOTE

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score

from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn import svm

from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
```

```
In [2]: import warnings
warnings.filterwarnings('ignore')
```

```
In [ ]: # Loan_ID : Unique Loan ID

# Gender : Male/ Female

# Married : Applicant married (Y/N)

# Dependents : Number of dependents

# Education : Applicant Education (Graduate/ Under Graduate)
```

```

# Self_Employed : Self employed (Y/N)

# ApplicantIncome : Applicant income

# CoapplicantIncome : Coapplicant income

# LoanAmount : Loan amount in thousands of dollars

# Loan_Amount_Term : Term of loan in months

# Credit_History : Credit history meets guidelines yes or no

# Property_Area : Urban/ Semi Urban/ Rural

# Loan_Status : Loan approved (Y/N) this is the target variable

```

```
In [13]: dataset = pd.read_csv('loan_prediction.csv')
dataset.shape
```

```
Out[13]: (614, 13)
```

```
In [14]: dataset.head()
```

```
Out[14]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncc
0	LP001002	Male	No	0	Graduate	No	5849	
1	LP001003	Male	Yes	1	Graduate	No	4583	150
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	235
4	LP001008	Male	No	0	Graduate	No	6000	

```
In [15]: dataset.tail()
```

```
Out[15]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantInr
609	LP002978	Female	No	0	Graduate	No	2900	
610	LP002979	Male	Yes	3+	Graduate	No	4106	
611	LP002983	Male	Yes	1	Graduate	No	8072	
612	LP002984	Male	Yes	2	Graduate	No	7583	
613	LP002990	Female	No	0	Graduate	Yes	4583	

```
In [16]: print('Number of Rows:',dataset.shape[0])
print('Number of Columns:',dataset.shape[1])
```

```

Number of Rows: 614
Number of Columns: 13

```

```
In [17]: dataset.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Loan_ID               614 non-null   object
1   Gender                601 non-null   object
2   Married               611 non-null   object
3   Dependents            599 non-null   object
4   Education             614 non-null   object
5   Self_Employed         582 non-null   object
6   ApplicantIncome       614 non-null   int64
7   CoapplicantIncome     614 non-null   float64
8   LoanAmount            592 non-null   float64
9   Loan_Amount_Term      600 non-null   float64
10  Credit_History        564 non-null   float64
11  Property_Area         614 non-null   object
12  Loan_Status           614 non-null   object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB

```

```
In [18]: dataset.isnull().sum()
```

```

Out[18]: Loan_ID           0
Gender           13
Married          3
Dependents       15
Education        0
Self_Employed    32
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       22
Loan_Amount_Term 14
Credit_History   50
Property_Area     0
Loan_Status       0
dtype: int64

```

```
In [19]: dataset.duplicated().sum()
```

```
Out[19]: 0
```

```
In [20]: dataset = dataset.dropna()
```

```
In [21]: dataset = dataset.drop('Loan_ID',axis=1)
```

```
In [22]: dataset.describe()
```

```
Out[22]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
count	480.000000	480.000000	480.000000	480.000000	480.000000
mean	5364.231250	1581.093583	144.735417	342.050000	0.854167
std	5668.251251	2617.692267	80.508164	65.212401	0.353307
min	150.000000	0.000000	9.000000	36.000000	0.000000
25%	2898.750000	0.000000	100.000000	360.000000	1.000000
50%	3859.000000	1084.500000	128.000000	360.000000	1.000000
75%	5852.500000	2253.250000	170.000000	360.000000	1.000000
max	81000.000000	33837.000000	600.000000	480.000000	1.000000

```
In [23]: dataset.corr()
```

```
Out[23]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
ApplicantIncome	1.000000	-0.112588	0.495310	-0.010838	-0.056152
CoapplicantIncome	-0.112588	1.000000	0.190740	-0.005775	-0.008692
LoanAmount	0.495310	0.190740	1.000000	0.050867	-0.040773
Loan_Amount_Term	-0.010838	-0.005775	0.050867	1.000000	0.032937
Credit_History	-0.056152	-0.008692	-0.040773	0.032937	1.000000

```
In [24]: dataset.cov()
```

```
Out[24]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
ApplicantIncome	3.212907e+07	-1.670551e+06	226029.825404	-4006.195303	-112.452
CoapplicantIncome	-1.670551e+06	6.852313e+06	40197.560179	-985.773871	-8.038
LoanAmount	2.260298e+05	4.019756e+04	6481.564505	267.057098	-1.159
Loan_Amount_Term	-4.006195e+03	-9.857739e+02	267.057098	4252.657203	0.758
Credit_History	-1.124526e+02	-8.038516e+00	-1.159751	0.758873	0.124

```
In [25]: dataset['Loan_Status'].value_counts()
```

```
Out[25]:
```

Y	332
N	148

Name: Loan_Status, dtype: int64

```
In [26]: dataset['Gender'].value_counts()
```

```
Out[26]:
```

Male	394
Female	86

Name: Gender, dtype: int64

```
In [27]: dataset['Married'].value_counts()
```

```
Out[27]:
```

Yes	311
No	169

Name: Married, dtype: int64

```
In [28]: dataset['Education'].value_counts()
```

```
Out[28]: Graduate      383  
Not Graduate    97  
Name: Education, dtype: int64
```

```
In [29]: dataset['Self_Employed'].value_counts()
```

```
Out[29]: No      414  
Yes      66  
Name: Self_Employed, dtype: int64
```

```
In [30]: dataset['Property_Area'].value_counts()
```

```
Out[30]: Semiurban    191  
Urban      150  
Rural      139  
Name: Property_Area, dtype: int64
```

```
In [31]: dataset['Credit_History'].value_counts()
```

```
Out[31]: 1.0      410  
0.0      70  
Name: Credit_History, dtype: int64
```


```
In [32]: dataset['Dependents'] = dataset['Dependents'].replace(to_replace='3+', value='4')
```

```
In [33]: dataset['Dependents'].value_counts()
```


```
Out[33]: 0      274  
2       85  
1       80  
4       41  
Name: Dependents, dtype: int64
```

```
In [34]: dataset['Gender'] = dataset['Gender'].map({'Male':1, 'Female':0}).astype('int')  
dataset['Married'] = dataset['Married'].map({'Yes':1, 'No':0}).astype('int')  
dataset['Education'] = dataset['Education'].map({'Graduate':1, 'Not Graduate':0}).astype('int')  
dataset['Self_Employed'] = dataset['Self_Employed'].map({'Yes':1, 'No':0}).astype('int')  
dataset['Property_Area'] = dataset['Property_Area'].map({'Rural':0, 'Urban':1, 'Semiurban':2}).astype('int')  
dataset['Loan_Status'] = dataset['Loan_Status'].map({'Y':1, 'N':0}).astype('int')
```


```
In [35]: sns.distplot(dataset['Loan_Status'])  
plt.show()
```




```
In [36]: sns.scatterplot(dataset['Gender'],dataset['Loan_Status'])  
plt.show()
```




```
In [37]: sns.scatterplot(dataset['Married'],dataset['Loan_Status'])  
plt.show()
```




```
In [28]: sns.scatterplot(dataset['Self_Employed'],dataset['Loan_Status'])  
plt.show()
```



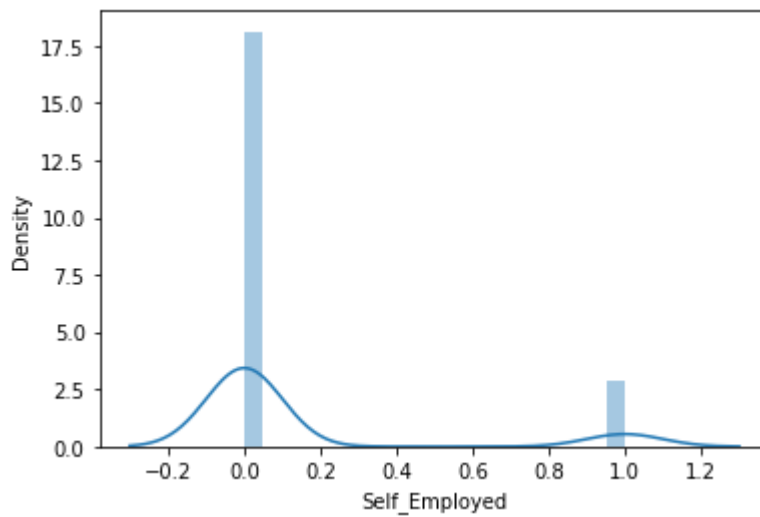
```
In [29]: sns.distplot(dataset['Gender'])  
plt.show()
```




```
In [30]: sns.distplot(dataset['Married'])  
plt.show()
```



```
In [31]: sns.distplot(dataset['Self_Employed'])  
plt.show()
```



```
In [32]: sns.distplot(dataset['Education'])
plt.show()
```



```
In [33]: dataset.dtypes
```

```
Out[33]: Gender          int32
Married          int32
Dependents       object
Education        int32
Self_Employed    int32
ApplicantIncome  int64
CoapplicantIncome float64
LoanAmount       float64
Loan_Amount_Term float64
Credit_History   float64
Property_Area    int32
Loan_Status      int32
dtype: object
```

```
In [34]: float_cols = [x for x in dataset.columns if dataset[x].dtypes=='float64']
```


```
In [35]: object_cols = [x for x in dataset.columns if dataset[x].dtypes=='object']
```

```
In [36]: int_cols = [x for x in dataset.columns if dataset[x].dtypes=='int64']
```




```
In [37]: plt.figure(figsize=(15,15))
plt.suptitle('Univariate Analysis of Numerical Features',fontweight='bold',font

for i in range(0,len(float_cols)):
    plt.subplot(10,4,i+1)
    sns.kdeplot(x=dataset[float_cols[i]],shade=True)
    plt.tight_layout()
```



```
In [38]: plt.figure(figsize=(20,20))
plt.suptitle('Univariate Analysis of Numerical Features',fontweight='bold',font

for i in range(0,len(float_cols)):
    plt.subplot(10,4,i+1)
    sns.boxplot(data=dataset,x=float_cols[i])
    plt.xlabel(float_cols[i])
    plt.tight_layout()
```




```
In [39]: def remove_outliers(in_dataset, in_cols):

    first_quartile = in_dataset[in_cols].quantile(0.25)
    third_quartile = in_dataset[in_cols].quantile(0.75)
    iqr = third_quartile - first_quartile
    upper_limit = third_quartile + 1.5 * iqr
    lower_limit = first_quartile - 1.5 * iqr
    in_dataset.loc[(in_dataset[in_cols] > upper_limit), in_cols] = upper_limit
    in_dataset.loc[(in_dataset[in_cols] < lower_limit), in_cols] = lower_limit
    return in_dataset
```

```
In [40]: for features in float_cols:
    dataset = remove_outliers(dataset,features)
```

```
In [41]: plt.figure(figsize=(20,20))
plt.suptitle('Univariate Analysis of Numerical Features',fontweight='bold',font

for i in range(0,len(float_cols)):
    plt.subplot(10,4,i+1)
    sns.boxplot(data=dataset,x=float_cols[i],color='red')
    plt.xlabel(float_cols[i])
    plt.tight_layout()
```



```
In [38]: dataset['Loan_Status'].value_counts()
```

```
Out[38]: 1    332  
0    148  
Name: Loan_Status, dtype: int64
```

```
In [39]: X = dataset.drop('Loan_Status',axis=1)  
y = dataset['Loan_Status']
```

```
In [46]: cols = ['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount', 'Loan_Amount_Term']
```

```
In [47]: st = StandardScaler()  
X[cols]=st.fit_transform(X[cols])
```

```
In [50]: from sklearn.model_selection import train_test_split  
from sklearn.model_selection import cross_val_score  
from sklearn.metrics import accuracy_score  
import numpy as np
```

```
In [51]: models_df = {}  
  
def model_val(model,X,y):  
    X_train,X_test,y_train,y_test=train_test_split(X,y,  
                                                    test_size=0.20,  
                                                    random_state=42)  
  
    model.fit(X_train,y_train)  
    y_pred=model.predict(X_test)  
    print(f"{model} accuracy is {accuracy_score(y_test,y_pred)}")  
  
    score = cross_val_score(model,X,y,cv=5)  
    print(f"{model} Avg cross val score is {np.mean(score)}")  
    models_df[model]=round(np.mean(score)*100,2)
```

```
In [58]: models_df
```

```
Out[58]: {LogisticRegression(): 80.21,  
SVC(): 79.79,  
DecisionTreeClassifier(): 72.5,  
RandomForestClassifier(): 80.0,  
GradientBoostingClassifier(): 78.54}
```

```
In [53]: from sklearn.linear_model import LogisticRegression  
model = LogisticRegression()  
model_val(model,X,y)
```

```
LogisticRegression() accuracy is 0.8229166666666666  
LogisticRegression() Avg cross val score is 0.8020833333333334
```

```
In [54]: from sklearn import svm  
model = svm.SVC()  
model_val(model,X,y)
```

```
SVC() accuracy is 0.8020833333333334  
SVC() Avg cross val score is 0.7979166666666667
```

```
In [55]: from sklearn.tree import DecisionTreeClassifier  
model = DecisionTreeClassifier()  
model_val(model,X,y)
```

DecisionTreeClassifier() accuracy is 0.7395833333333334
DecisionTreeClassifier() Avg cross val score is 0.725

```
In [56]: from sklearn.ensemble import RandomForestClassifier
model =RandomForestClassifier()
model_val(model,X,y)
```

RandomForestClassifier() accuracy is 0.8125
RandomForestClassifier() Avg cross val score is 0.8

```
In [57]: from sklearn.ensemble import GradientBoostingClassifier
model =GradientBoostingClassifier()
model_val(model,X,y)
```

GradientBoostingClassifier() accuracy is 0.7604166666666666
GradientBoostingClassifier() Avg cross val score is 0.7854166666666667

```
In [59]: log_reg_grid={"C":np.logspace(-4,4,20),
                      "solver":["liblinear"]}
```

```
In [60]: rs_log_reg=RandomizedSearchCV(LogisticRegression(),
                                     param_distributions=log_reg_grid,
                                     n_iter=20,cv=5,verbose=True)
```

```
In [61]: rs_log_reg.fit(X,y)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

```
Out[61]: RandomizedSearchCV
  estimator: LogisticRegression
    LogisticRegression
```

```
In [62]: rs_log_reg.best_score_
```

```
Out[62]: 0.8020833333333334
```

```
In [63]: rs_log_reg.best_params_
```

```
Out[63]: {'solver': 'liblinear', 'C': 0.23357214690901212}
```

```
In [64]: svc_grid = {'C':[0.25,0.50,0.75,1],"kernel":["linear"]}
```

```
In [65]: rs_svc=RandomizedSearchCV(svm.SVC(),
                                   param_distributions=svc_grid,
                                   cv=5,
                                   n_iter=20,
                                   verbose=True)
```

```
In [66]: rs_svc.fit(X,y)
```

Fitting 5 folds for each of 4 candidates, totalling 20 fits

```
Out[66]: ▸ RandomizedSearchCV
          ▸ estimator: SVC
            ▸ SVC
```

```
In [67]: rs_svc.best_score_
```

```
Out[67]: 0.8083333333333333
```

```
In [68]: rs_svc.best_params_
```

```
Out[68]: {'kernel': 'linear', 'C': 0.25}
```

```
In [69]: RandomForestClassifier()
```

```
Out[69]: ▾ RandomForestClassifier
          RandomForestClassifier()
```

```
In [70]: rf_grid={'n_estimators':np.arange(10,1000,10),
                  'max_features':['auto','sqrt'],
                  'max_depth':[None,3,5,10,20,30],
                  'min_samples_split':[2,5,20,50,100],
                  'min_samples_leaf':[1,2,5,10]}
}
```

```
In [71]: rs_rf=RandomizedSearchCV(RandomForestClassifier(),
                                   param_distributions=rf_grid,
                                   cv=5,
                                   n_iter=20,
                                   verbose=True)
```

```
In [72]: rs_rf.fit(X,y)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

```
Out[72]: ▸ RandomizedSearchCV
          ▸ estimator: RandomForestClassifier
            ▸ RandomForestClassifier
```

```
In [73]: rs_rf.best_score_
```

```
Out[73]: 0.8083333333333333
```

```
In [74]: rs_rf.best_params_
```

```
Out[74]: {'n_estimators': 780,
          'min_samples_split': 20,
          'min_samples_leaf': 1,
          'max_features': 'auto',
          'max_depth': 5}
```

```
In [78]: X_res,y_res = SMOTE().fit_resample(X,y)
```

```
In [79]: X_train,X_test,y_train,y_test = train_test_split(X_res,y_res,test_size=0.20,rai
```

```
In [80]: st = StandardScaler()  
X_train = st.fit_transform(X_train)  
X_test = st.fit_transform(X_test)
```

```
In [81]: model_df = {}  
  
def model_val(model,X,y):  
    X_train,X_test,y_train,y_test = train_test_split(X_res,y_res,test_size=0.20,  
    model.fit(X_train,y_train)  
    y_pred = model.predict(X_test)  
    print(f'{model} Accuracy is {accuracy_score(y_test,y_pred)}')  
  
    score = cross_val_score(model,X,y,cv=5,n_jobs=-1)  
    print(f'{model} Average cross val score is {np.mean(score)}')  
    model_df[model] = round(np.mean(score)*100,2)
```

```
In [82]: model = LogisticRegression()  
model_val(model,X,y)  
  
LogisticRegression() Accuracy is 0.7593984962406015  
LogisticRegression() Average cross val score is 0.8020833333333334
```

```
In [49]: model = DecisionTreeClassifier()  
model_val(model,X,y)  
  
DecisionTreeClassifier() Accuracy is 0.706766917293233  
DecisionTreeClassifier() Average cross val score is 0.5729166666666667
```

```
In [50]: model = RandomForestClassifier()  
model_val(model,X,y)  
  
RandomForestClassifier() Accuracy is 0.7368421052631579  
RandomForestClassifier() Average cross val score is 0.6375
```

```
In [51]: model = GradientBoostingClassifier()  
model_val(model,X,y)  
  
GradientBoostingClassifier() Accuracy is 0.7593984962406015  
GradientBoostingClassifier() Average cross val score is 0.675
```

```
In [52]: model_df
```

```
Out[52]: {LogisticRegression(): 69.17,  
DecisionTreeClassifier(): 57.29,  
RandomForestClassifier(): 63.75,  
GradientBoostingClassifier(): 67.5}
```

```
In [53]: C = 1.0  
svm = svm.SVC(kernel='linear',C=C)  
svm.fit(X_train,y_train)
```

```
Out[53]: ▼ SVC  
SVC(kernel='linear')
```

```
In [54]: y_pred = svm.predict(X_test)
```

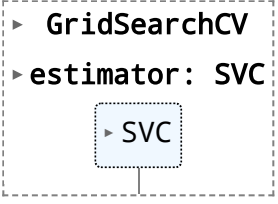
```
In [55]: accuracy_score(y_test,y_pred)
```

```
Out[55]: 0.7293233082706767
```

```
In [56]: parameters = [  
    {'kernel':['linear'],'C': [1,10]},  
    {'kernel':['poly'],'C':[1, 10]},  
    {'kernel':['rbf'],'gamma': [1e-3,1e-4],'C': [1,10]}  
]
```

```
In [57]: gr = GridSearchCV(svm,parameters,n_jobs=-1)  
gr.fit(X_train,y_train)
```

```
Out[57]:
```



```
  ▸ GridSearchCV  
  ▸ estimator: SVC  
    ▸ SVC
```

```
In [58]: y_pred_ = gr.predict(X_test)
```

```
In [59]: accuracy_score(y_test,y_pred_)
```

```
Out[59]: 0.7293233082706767
```

```
In [60]: from sklearn.model_selection import StratifiedKFold  
from sklearn.linear_model import LogisticRegression  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.calibration import CalibratedClassifierCV  
from sklearn.neural_network import MLPClassifier  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis  
from sklearn.naive_bayes import GaussianNB  
from sklearn.ensemble import AdaBoostClassifier  
from sklearn.ensemble import GradientBoostingClassifier  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.ensemble import ExtraTreesClassifier  
from sklearn.linear_model import SGDClassifier  
from sklearn.svm import SVC  
from xgboost import XGBClassifier
```

```
In [61]: estimators = [  
    ('rf', RandomForestClassifier(n_estimators = 10, random_state = 42)),  
    ('knn', KNeighborsClassifier(n_neighbors = 10)),  
    ('gbc', GradientBoostingClassifier()),  
    ('lr', LogisticRegression()),  
    ('ccv', CalibratedClassifierCV()),  
    ('mlp', MLPClassifier()),  
    ('dt', DecisionTreeClassifier()),  
    ('lda', LinearDiscriminantAnalysis()),  
    ('gnb', GaussianNB()),
```

```

        ('adb', AdaBoostClassifier()),
        ('etc', ExtraTreesClassifier()),
        ('sgd', SGDClassifier()),
        ('svm', SVC()),
        ('xgb', XGBClassifier(n_estimators= 10, random_state = 42))
    ]

```

In [62]: **from** sklearn.ensemble **import** StackingClassifier


```

st = StackingClassifier(
    estimators = estimators,
    final_estimator = LogisticRegression(),
    cv = 10,
    stack_method='predict',
    n_jobs=-1
)

```

In [63]: st.fit(X_train,y_train)

Out[63]:



In [64]: y_pred = st.predict(X_test)

In [65]: accuracy_score(y_test,y_pred)

Out[65]: 0.6992481203007519

```

In [ ]: from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.model_selection import GridSearchCV
from xgboost import XGBClassifier

```

```

In [66]: dt_parameters = {
    'criterion' : ['gini', 'entropy', 'log_loss'],
    'splitter' : ['best', 'random'],
    'max_depth' : range(1,10,1),
    'min_samples_split' : range(2,10,2),
    'min_samples_leaf' : range(1,5,1),
    'max_features' : ['auto', 'sqrt', 'log2']
}

bg_parameters = {
    'n_estimators' : [5, 10, 15],
    'max_samples' : range(2, 10, 1),
    'max_features' : range(2, 10, 3)
}

```



```
In [ ]: rs_rf.fit(X,y)
```

```
In [ ]: rs_rf.best_score_
```

```
In [68]: import sklearn
from sklearn.model_selection import train_test_split, GridSearchCV, ShuffleSplit
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.naive_bayes import BernoulliNB, GaussianNB
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

model_list = [('lr', LogisticRegression()), ('dt', DecisionTreeClassifier()), ('rf',
                                         RandomForestClassifier()), ('svm', SVC()), ('bnb', BernoulliNB()), ('gnb', GaussianNB())]

for model_name, model in model_list:
    m = model.fit(X_train, y_train)
    y_pred = model.predict(X_train)
    print(f'{model_name} : {accuracy_score(y_train, y_pred)}')
```

```
lr : 0.6610169491525424
dt : 1.0
rf : 0.9887005649717514
adb : 0.7664783427495292
svm : 0.7532956685499058
bnb : 0.664783427495292
gnb : 0.6892655367231638
```

```
In [69]: for model_name, model in model_list:
    m = model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    print(f'{model_name} : {accuracy_score(y_test, y_pred)}')
```

```
lr : 0.6917293233082706
dt : 0.6541353383458647
rf : 0.7368421052631579
adb : 0.5939849624060151
svm : 0.6691729323308271
bnb : 0.6992481203007519
gnb : 0.7142857142857143
```

```
In [ ]: algos = {
    'rf': {
        'model': RandomForestClassifier(),
        'params': {
            'n_estimators': [10, 20, 30, 40],
            'criterion': ['gini', 'entropy'],
            'max_depth': [10, 20, 30],
            'min_samples_split': [2, 4, 6]
        }
    },
    'dt': {
        'model': DecisionTreeClassifier(),
        'params': {
            'criterion': ['gini', 'entropy'],
            'max_depth': [10, 20, 30],
            'splitter': ['best', 'random'],
        }
    }
}
```

```

    },
    'lr':{
        'model':LogisticRegression(),
        'params':{
            'penalty':['l1','l2'],
            'C':[0.1,0.01,1,0.5],
            'solver':['liblinear','lbfgs']
        }
    },
    'svm':{
        'model':SVC(),
        'params':{
            'C':[0.1,0.01,1,0.5],
            'kernel':['linear','poly','rbf']
        }
    },
    'adb':{
        'model':AdaBoostClassifier(),
        'params':{
            'n_estimators':[10,20,30,50],
            'learning_rate':[0.1,1,0.01]
        }
    }
}

scores = []

cv = ShuffleSplit(n_splits=10,test_size=0.2,random_state=42)

for model_name,config in algos.items():
    gd = GridSearchCV(estimator=config['model'],param_grid=config['params'],cv=cv)
    gd.fit(X,y)
    scores.append({'model_name':model_name,'best_score':gd.best_score_,'best_p

scores = pd.DataFrame(scores)
scores.head()

```

```
In [84]: X = dataset.drop('Loan_Status',axis=1)
y = dataset['Loan_Status']
```

```
In [85]:
```

```
In [86]: from sklearn.preprocessing import StandardScaler
st = StandardScaler()
X[cols]=st.fit_transform(X[cols])
```

```
In [87]: from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score
import numpy as np
```

```
In [88]: model_df = {}

def model_val(model,X,y):
    X_train,X_test,y_train,y_test=train_test_split(X,y,
                                                    test_size=0.20,
                                                    random_state=42)

    model.fit(X_train,y_train)
    y_pred=model.predict(X_test)
    print(f"{model} accuracy is {accuracy_score(y_test,y_pred)}")

    score = cross_val_score(model,X,y,cv=5)
    print(f"{model} Avg cross val score is {np.mean(score)}")
    model_df[model]=round(np.mean(score)*100,2)
```

```
In [85]: model_df
```

```
Out[85]: {LogisticRegression(): 80.21}
```

```
In [84]: X = dataset.drop('Loan_Status',axis=1)
y = dataset['Loan_Status']
```

```
In [86]: rf = RandomForestClassifier(n_estimators=270,
    min_samples_split=5,
    min_samples_leaf=5,
    max_features='sqrt',
    max_depth=5)
```

```
In [87]: rf.fit(X,y)
```

```
Out[87]: ▼ RandomForestClassifier
RandomForestClassifier(max_depth=5, min_samples_leaf=5, min_samples_split=5,
                        n_estimators=270)
```

```
In [88]: import pickle
import joblib
```

```
In [92]: pickle.dump(rf,open('loan_status_predict.pkl','wb'))
```

```
In [91]: model = pickle.load(open('loan_status_predict.pkl','rb'))
```

```
In [93]: df = pd.DataFrame({
    'Gender':1,
    'Married':1,
    'Dependents':2,
    'Education':0,
    'Self_Employed':0,
    'ApplicantIncome':2889,
    'CoapplicantIncome':0.0,
    'LoanAmount':45,
    'Loan_Amount_Term':180,
    'Credit_History':0,
    'Property_Area':1
},index=[0])
```

```
In [94]: df
```

```
Out[94]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanA
0	1	1	2	0	0	2889	0.0	

```
In [95]: result = model.predict(df)
```

```
In [97]: if result==1:
    print('Loan Approved')
else:
    print('Loan Not Approved')
```

Loan Not Approved

```
In [104... p = model.predict(df)

prob = model.predict_proba(df)
if p == 1:
    print('Loan Approved!')
    print(f'You will be Loan Approved! with Probability of {prob[0][1]:.2f}')
else:
    print('Not-Loan Approved!')
```

Not-Loan Approved!

```
In [98]: from tkinter import *
```

```
In [107... def show_entry():

    p1 = float(e1.get())
    p2 = float(e2.get())
    p3 = float(e3.get())
    p4 = float(e4.get())
    p5 = float(e5.get())
    p6 = float(e6.get())
    p7 = float(e7.get())
    p8 = float(e8.get())
    p9 = float(e9.get())
    p10 = float(e10.get())
    p11 = float(e11.get())

    model = pickle.load(open('loan_status_predict.pkl','rb'))
    df = pd.DataFrame({
```

```

        'Gender':p1,
        'Married':p2,
        'Dependents':p3,
        'Education':p4,
        'Self_Employed':p5,
        'ApplicantIncome':p6,
        'CoapplicantIncome':p7,
        'LoanAmount':p8,
        'Loan_Amount_Term':p9,
        'Credit_History':p10,
        'Property_Area':p11
    },index=[0])
    result = model.predict(df)

    if result == 1:
        Label(master, text="Loan Approved").grid(row=31)
    else:
        Label(master, text="Loan Not Approved").grid(row=31)

master =Tk()
master.title("Loan Status Prediction Using Machine Learning")
label = Label(master,text = "Loan Status Prediction",bg = "red",
               fg = "white").grid(row=0,columnspan=2)

Label(master,text = "Gender [1:Male ,0:Female]").grid(row=1)
Label(master,text = "Married [1:Yes,0:No]").grid(row=2)
Label(master,text = "Dependents [1,2,3,4]").grid(row=3)
Label(master,text = "Education").grid(row=4)
Label(master,text = "Self_Employed").grid(row=5)
Label(master,text = "ApplicantIncome").grid(row=6)
Label(master,text = "CoapplicantIncome").grid(row=7)
Label(master,text = "LoanAmount").grid(row=8)
Label(master,text = "Loan_Amount_Term").grid(row=9)
Label(master,text = "Credit_History").grid(row=10)
Label(master,text = "Property_Area").grid(row=11)


e1 = Entry(master)
e2 = Entry(master)
e3 = Entry(master)
e4 = Entry(master)
e5 = Entry(master)
e6 = Entry(master)
e7 = Entry(master)
e8 = Entry(master)
e9 = Entry(master)
e10 = Entry(master)
e11 = Entry(master)


e1.grid(row=1,column=1)
e2.grid(row=2,column=1)
e3.grid(row=3,column=1)
e4.grid(row=4,column=1)
e5.grid(row=5,column=1)
e6.grid(row=6,column=1)
e7.grid(row=7,column=1)
e8.grid(row=8,column=1)
e9.grid(row=9,column=1)

```

```
e10.grid(row=10,column=1)
e11.grid(row=11,column=1)
```

```
Button(master,text="Predict",command=show_entry).grid()
```

```
mainloop()
```

Loan Status Prediction Using Machine Learning

Loan Status Prediction	
Gender [1:Male,0:Female]	<input type="text" value="1"/>
Married [1:Yes,0:No]	<input type="text" value="1"/>
Dependents [1,2,3,4]	<input type="text" value="2"/>
Education	<input type="text" value="0"/>
Self_Employed	<input type="text" value="0"/>
ApplicantIncome	<input type="text" value="2889"/>
CoapplicantIncome	<input type="text" value="0.0"/>
LoanAmount	<input type="text" value="45"/>
Loan_Amount_Term	<input type="text" value="180"/>
Credit_History	<input type="text" value="0"/>
Property_Area	<input type="text" value="1"/>
<input type="button" value="Predict"/>	
Loan Not Approved	

CONNECT WITH ME:

[LinkedIn](#) [GitHub](#) [kaggle](#) [Medium](#)

PRASADMJADHAV2